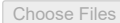


```
from google.colab import files
uploaded = files.upload()
```

 No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.  
Saving HR-Employee-Attrition.csv to HR-Employee-Attrition.csv

```
import pandas as pd
```

```
df = pd.read_csv("HR-Employee-Attrition.csv") # Use the exact uploaded filename
```

```
# Basic info
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   1470 non-null   int64
1   Attrition                            1470 non-null   object
2   BusinessTravel                       1470 non-null   object
3   DailyRate                           1470 non-null   int64
4   Department                           1470 non-null   object
5   DistanceFromHome                    1470 non-null   int64
6   Education                           1470 non-null   int64
7   EducationField                       1470 non-null   object
8   EmployeeCount                       1470 non-null   int64
9   EmployeeNumber                      1470 non-null   int64
10  EnvironmentSatisfaction              1470 non-null   int64
11  Gender                              1470 non-null   object
12  HourlyRate                          1470 non-null   int64
13  JobInvolvement                      1470 non-null   int64
14  JobLevel                            1470 non-null   int64
15  JobRole                             1470 non-null   object
16  JobSatisfaction                     1470 non-null   int64
17  MaritalStatus                       1470 non-null   object
18  MonthlyIncome                       1470 non-null   int64
19  MonthlyRate                         1470 non-null   int64
20  NumCompaniesWorked                  1470 non-null   int64
21  Over18                              1470 non-null   object
22  OverTime                            1470 non-null   object
23  PercentSalaryHike                   1470 non-null   int64
24  PerformanceRating                   1470 non-null   int64
25  RelationshipSatisfaction             1470 non-null   int64
26  StandardHours                       1470 non-null   int64
27  StockOptionLevel                    1470 non-null   int64
28  TotalWorkingYears                   1470 non-null   int64
29  TrainingTimesLastYear               1470 non-null   int64
30  WorkLifeBalance                     1470 non-null   int64
31  YearsAtCompany                      1470 non-null   int64
32  YearsInCurrentRole                  1470 non-null   int64
33  YearsSinceLastPromotion              1470 non-null   int64
34  YearsWithCurrManager                 1470 non-null   int64
dtypes: int64(26), object(9)
memory usage: 402.1+ KB
```

```
df.head()
```


```
<img alt="Table view icon" data-bbox="70 725 90 735"/>
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Other	1	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	1	
4	27	No	Travel_Rarely	591	Research & Development	2	1	Medical	1	

5 rows × 35 columns

```
import seaborn as sns
import matplotlib.pyplot as plt
```

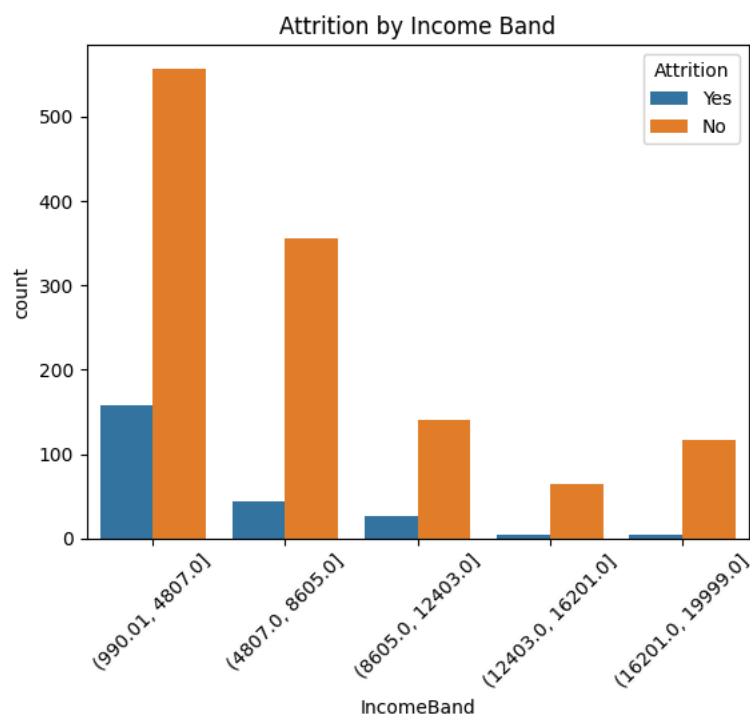
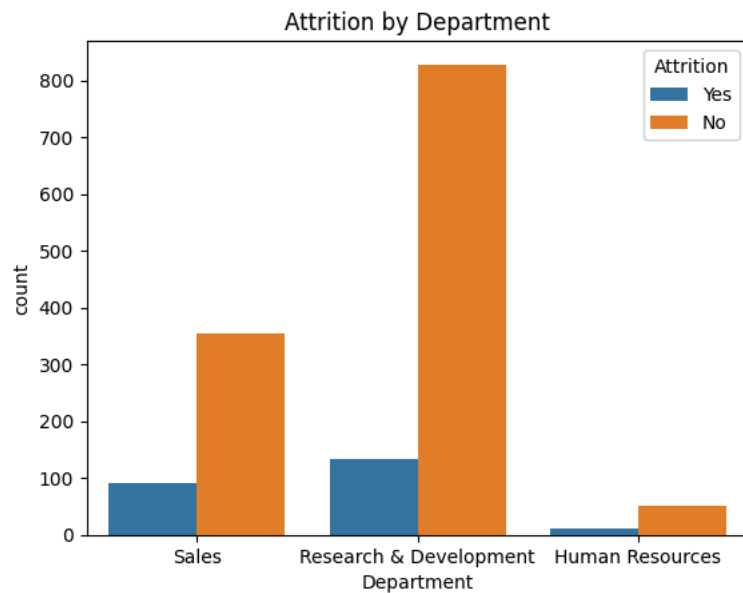
```
# Attrition Count
sns.countplot(x='Attrition', data=df)
plt.title('Overall Attrition Count')
plt.show()
```



```
# Attrition by Department
sns.countplot(x='Department', hue='Attrition', data=df)
plt.title('Attrition by Department')
plt.show()

# Income Bands
df['IncomeBand'] = pd.cut(df['MonthlyIncome'], bins=5)
sns.countplot(x='IncomeBand', hue='Attrition', data=df)
plt.xticks(rotation=45)
plt.title('Attrition by Income Band')
plt.show()

# Promotion vs Attrition
sns.countplot(x='YearsSinceLastPromotion', hue='Attrition', data=df)
plt.title('Attrition vs. Years Since Last Promotion')
plt.show()
```



```
# Create a copy for modeling
df_model = df.copy()

# Drop the 'IncomeBand' column if it exists
if 'IncomeBand' in df_model.columns:
    df_model.drop('IncomeBand', axis=1, inplace=True)

# Label encode categorical features
le = LabelEncoder()
for col in df_model.select_dtypes(include='object'):
    df_model[col] = le.fit_transform(df_model[col])

# Split features and target
X = df_model.drop(['Attrition'], axis=1)
y = df_model['Attrition']

# Train-test split
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train model
from sklearn.tree import DecisionTreeClassifier
model = DecisionTreeClassifier()
model.fit(X_train, y_train)

# Evaluate model
from sklearn.metrics import classification_report, confusion_matrix
y_pred = model.predict(X_test)
```

```
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

↗ Confusion Matrix:

```
[[219  36]
 [ 30   9]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.88	0.86	0.87	255
1	0.20	0.23	0.21	39
accuracy			0.78	294
macro avg	0.54	0.54	0.54	294
weighted avg	0.79	0.78	0.78	294

```
df['Attrition'].value_counts()
```

↗

Attrition	count
No	1233
Yes	237

dtype: int64

```
from imblearn.over_sampling import SMOTE
```

```
smote = SMOTE(random_state=42)
X_sm, y_sm = smote.fit_resample(X, y)
```

```
# New shape after balancing
print(y_sm.value_counts())
```

↗ Attrition

```
1    1233
0    1233
Name: count, dtype: int64
```

```
X_train, X_test, y_train, y_test = train_test_split(X_sm, y_sm, test_size=0.2, random_state=42)
```

```
model = DecisionTreeClassifier()
model.fit(X_train, y_train)
```

```
y_pred = model.predict(X_test)
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

↗ Confusion Matrix:

```
[[181  69]
 [ 38 206]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.72	0.77	250
1	0.75	0.84	0.79	244
accuracy			0.78	494
macro avg	0.79	0.78	0.78	494
weighted avg	0.79	0.78	0.78	494

```
from sklearn.linear_model import LogisticRegression
```

```
log_model = LogisticRegression(max_iter=2000)
log_model.fit(X_train, y_train)
```

```
y_pred_log = log_model.predict(X_test)
print("Logistic Regression Report:\n", classification_report(y_test, y_pred_log))
```

↗ Logistic Regression Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.91	0.98	0.94	255
1	0.68	0.33	0.45	39
accuracy			0.89	294
macro avg	0.79	0.65	0.69	294
weighted avg	0.88	0.89	0.87	294

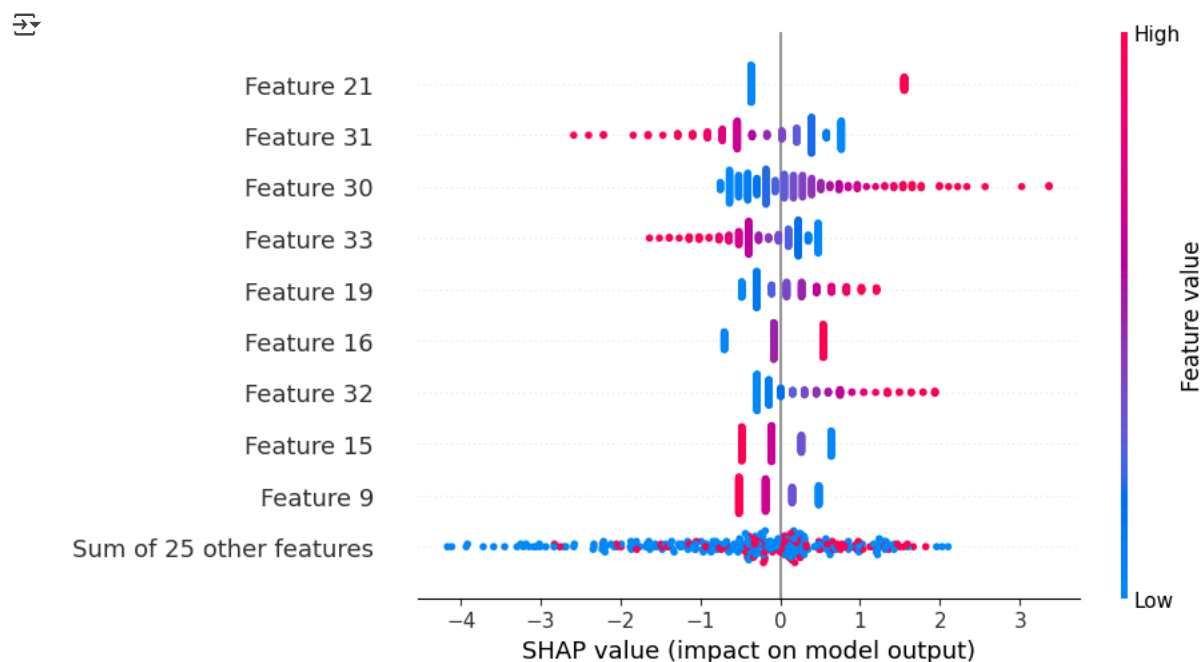
```
!pip install shap
```

```
Requirement already satisfied: shap in /usr/local/lib/python3.11/dist-packages (0.47.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from shap) (2.0.2)
Requirement already satisfied: scipy in /usr/local/lib/python3.11/dist-packages (from shap) (1.15.2)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.11/dist-packages (from shap) (1.6.1)
Requirement already satisfied: pandas in /usr/local/lib/python3.11/dist-packages (from shap) (2.2.2)
Requirement already satisfied: tqdm>=4.27.0 in /usr/local/lib/python3.11/dist-packages (from shap) (4.67.1)
Requirement already satisfied: packaging>20.9 in /usr/local/lib/python3.11/dist-packages (from shap) (24.2)
Requirement already satisfied: slicer==0.0.8 in /usr/local/lib/python3.11/dist-packages (from shap) (0.0.8)
Requirement already satisfied: numba>=0.54 in /usr/local/lib/python3.11/dist-packages (from shap) (0.60.0)
Requirement already satisfied: cloudpickle in /usr/local/lib/python3.11/dist-packages (from shap) (3.1.1)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.11/dist-packages (from shap) (4.13.2)
Requirement already satisfied: llvmlite<0.44,>=0.43.0dev0 in /usr/local/lib/python3.11/dist-packages (from numba>=0.54->shap) (0.43)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas->shap) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas->shap) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas->shap) (2025.2)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn->shap) (1.4.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn->shap) (3.6.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.8.2->pandas->shap) (1.17)
```

```
import shap
import matplotlib.pyplot as plt

# For tree models, use TreeExplainer; for linear models, use LinearExplainer
explainer = shap.Explainer(log_model, X_train)
shap_values = explainer(X_test)

# Summary plot (feature importance)
shap.plots.beeswarm(shap_values)
```



```
# Display SHAP value for a single prediction
shap.initjs()
shap.plots.force(shap_values[0])
```